

Linear Regression:

```
import numpy as np

class LinearRegression:
    def __init__(self, lr=0.001, n_iters=1000):
        self.lr = lr
        self.n_iters = n_iters
        self.w = None
        self.b = None

    def fit(self, X, y):
        n_samples, n_features = X.shape
        self.w = np.zeros(n_features)
        self.b = 0

        for _ in range(self.n_iters):
            y_pred = np.dot(X, self.w) + self.b

            dw = (1/n_samples) * np.dot(X.T, (y_pred - y))
            db = (1/n_samples) * np.sum(y_pred - y)

            self.w = self.w - self.lr * dw
            self.b = self.b - self.lr * db

    def predict(self, X):
        y_pred = np.dot(X, self.w) + self.b
        return y_pred
```

Logistic Regression:

```
import numpy as np

def sigmoid(x):
    return 1/(1+np.exp(-x))

class LogisticRegression():

    def __init__(self, lr=0.001, n_iters=1000):
        self.lr = lr
        self.n_iters = n_iters
        self.w = None
        self.b = None

    def fit(self, X, y):
        n_samples, n_features = X.shape
        self.w = np.zeros(n_features)
        self.b = 0.0

        for i in range(self.n_iters):
            linear_pred = np.dot(X, self.w) + self.b
            predictions = sigmoid(linear_pred)

            dw = (1/n_samples) * np.dot(X.T, (predictions - y))
            db = (1/n_samples) * np.sum(predictions-y)

            self.w = self.w - self.lr*dw
            self.b = self.b - self.lr*db

    def predict(self, X):
        linear_pred = np.dot(X, self.w) + self.b
        y_pred = sigmoid(linear_pred)
        class_pred = [0 if y<=0.5 else 1 for y in y_pred]
        return class_pred
```

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from LinearRegression import LinearRegression
from sklearn.metrics import mean_squared_error

# load dataset
df = pd.read_csv("dataset/advertising/advertising.csv")

# i. Understand the Dataset & cleanup (if required).
print("\nSample Data: \n", df.head())

# data info
print("Data Info: \n")
df.info()

# describe data
print("\nData Description: \n", df.describe())

# check for missing data
print("\nMissing Values: \n", df.isnull().sum())

# if data consists missing data
df = df.dropna(subset=["TV Ad Budget ($)", "Radio Ad Budget ($)",
"Newspaper Ad Budget ($)", "Sales ($)"])

# Visualize data
plt.figure(figsize=(10, 6))
sns.scatterplot(data=df, x="Radio Ad Budget ($)", y="Sales ($)")
plt.show()

# sales w.r.t Radio features.
X1 = df[['Radio Ad Budget ($)']]
y1 = df['Sales ($)']

# sales w.r.t attribute tv.
X2 = df[['TV Ad Budget ($)']]
y2 = df['Sales ($)']

```

```

# sales w.r.t attribute newspaper.
X3 = df[['Newspaper Ad Budget ($)']]
y3 = df['Sales ($)']

# sales w.r.t Radio and TV
X4 = df[['Radio Ad Budget ($)', 'TV Ad Budget ($)']]
y4 = df['Sales ($)']

# sales w.r.t Newspaper and TV
X5 = df[['Newspaper Ad Budget ($)', 'TV Ad Budget ($)']]
y5 = df['Sales ($)']

# sales w.r.t Newspaper and Radio
X6 = df[['Newspaper Ad Budget ($)', 'Radio Ad Budget ($)']]
y6 = df['Sales ($)']

X_train, X_test, y_train, y_test = train_test_split(X1, y1,
test_size=0.2, random_state=42)

# use only for tv column
# X_train = X_train/10
# X_test = X_test/10

model = LinearRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

# iii. Also evaluate the model using scores RMSE

rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print(f"Root Mean Squared Error (RMSE): {rmse}")

# 1-3
plt.scatter(X_test, y_test, color='blue', label='Actual')
plt.plot(X_test, y_pred, color='red', label='Predicted')
plt.title('Actual vs Predicted Sales')
plt.xlabel('Radio')
plt.ylabel('Sales')
plt.legend()
plt.show()

```

```
# 4-6
plt.figure(figsize=(12,6))

indices = np.arange(len(y_test))
plt.bar(indices - 0.2, y_test, width=0.4, label='Actual Sales',
color='blue')
plt.bar(indices + 0.2, y_pred, width=0.4, label='Predicted Sales',
color='red')

plt.title('Comparison of Actual vs Predicted Sales')
plt.xlabel('Data Points')
plt.ylabel('Sales')
plt.legend()

plt.show()
```




```
Sample Data:
   Unnamed: 0  TV Ad Budget ($)  Radio Ad Budget ($)  Newspaper Ad Budget ($)  Sales ($)
0           1           230.1           37.8           69.2           22.1
1           2            44.5           39.3           45.1           10.4
2           3            17.2           45.9           69.3            9.3
3           4           151.5           41.3           58.5           18.5
4           5           180.8           10.8           58.4           12.9

Data Info:

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Unnamed: 0                            200 non-null    int64
 1   TV Ad Budget ($)                      200 non-null    float64
 2   Radio Ad Budget ($)                   200 non-null    float64
 3   Newspaper Ad Budget ($)                200 non-null    float64
 4   Sales ($)                             200 non-null    float64
dtypes: float64(4), int64(1)
```

```
Data Description:
   Unnamed: 0  TV Ad Budget ($)  Radio Ad Budget ($)  Newspaper Ad Budget ($)  Sales ($)
count  200.000000      200.000000      200.000000      200.000000  200.000000
mean   100.500000      147.042500       23.264000       30.554000   14.022500
std     57.879185       85.854236       14.846809       21.778621    5.217457
min      1.000000        0.700000        0.000000        0.300000    1.600000
25%     50.750000       74.375000        9.975000       12.750000   10.375000
50%    100.500000      149.750000       22.900000       25.750000   12.900000
75%    150.250000      218.825000       36.525000       45.100000   17.400000
max    200.000000      296.400000       49.600000      114.000000   27.000000

Missing Values:
   Unnamed: 0      0
TV Ad Budget ($)      0
Radio Ad Budget ($)    0
Newspaper Ad Budget ($) 0
Sales ($)             0
dtype: int64
Root Mean Squared Error (RMSE): 5.845791284816158
```



```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from LinearRegression import LinearRegression
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv("dataset/car/car data.csv")

print(df.head())
print(df.info())

# Convert 'transmission' column to binary values: 0 for 'Manual', 1 for
'Automatic'
df["Transmission"] = df["Transmission"].map({"Manual": 0, "Automatic":
1})
df["Fuel_Type"] = df["Fuel_Type"].map({"Petrol": 0, "Diesel": 1})
df["Seller_Type"] = df["Seller_Type"].map({"Dealer": 0, "Individual":
1})
df["Year"] = df["Year"] / 100
df["Kms_Driven"] = df["Kms_Driven"] / 1000
df = df.dropna()

sns.scatterplot(x="Kms_Driven", y="Selling_Price", data=df)
plt.title("Kms_Driven vs Selling_Price")
plt.show()

# 7. Selling prices w.r.t year bought
X1 = df[["Year"]] # Input
y1 = df["Selling_Price"] # Output

# 8. Selling prices w.r.t km driven
X2 = df[["Kms_Driven"]] # Input
y2 = df["Selling_Price"] # Output

# 9. Selling prices w.r.t transmission
X3 = df[["Transmission"]] # Input
y3 = df["Selling_Price"] # Output

# 10. Selling prices w.r.t owner

```

```

X4 = df[["Owner"]] # Input
y4 = df["Selling_Price"] # Output

# 11. Selling prices w.r.t year bought and km driven
X5 = df[["Year", "Kms_Driven"]] # Input
y5 = df["Selling_Price"] # Output

# 12. Selling prices w.r.t year bought and transmission
X6 = df[["Year", "Transmission"]] # Input
y6 = df["Selling_Price"] # Output

# 13-14. Selling prices w.r.t year bought and owner
X7 = df[["Year", "Owner"]] # Input
y7 = df["Selling_Price"] # Output

# 15. Selling prices w.r.t km driven and transmission
X8 = df[["Kms_Driven", "Transmission"]] # Input
y8 = df["Selling_Price"] # Output

# 16. Selling prices w.r.t km driven and owner
X9 = df[["Kms_Driven", "Owner"]] # Input
y9 = df["Selling_Price"] # Output

# 17-18. Selling prices w.r.t transmission and owner
X10 = df[["Transmission", "Owner"]] # Input
y10 = df["Selling_Price"] # Output

X11 = df.drop(["Selling_Price", "Car_Name"], axis=1)
y11 = df["Selling_Price"]

X_train, X_test, y_train, y_test = train_test_split(
    X8, y8, test_size=0.2, random_state=42
)

model = LinearRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
y_pred = np.clip(y_pred, a_min=0, a_max=None)

rmse = np.sqrt(np.mean((y_test - y_pred) ** 2))
print(f"Root Mean Squared Error (RMSE): {rmse}")

```






```

# 7-10
# plt.scatter(X_test, y_test, color='blue', label='Actual')
# plt.plot(X_test, y_pred, color='red', label='Predicted')
# plt.title('Year vs Selling_Price')
# plt.xlabel('Year')
# plt.ylabel('Selling_Price')
# plt.legend()
# plt.show()

# 11-18
plt.figure(figsize=(12,6))
indices = np.arange(len(y_test))
plt.bar(indices - 0.2, y_test, width=0.4, label='Actual Price',
color='blue')
plt.bar(indices + 0.2, y_pred, width=0.4, label='Selling Price',
color='red')
plt.title('Comparison of Actual vs Predicted Price')
plt.xlabel('Data Points')
plt.ylabel('Selling Price')
plt.legend()
plt.show()

```





```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as snb
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import confusion_matrix, accuracy_score,
recall_score, precision_score, f1_score

df = pd.read_csv("dataset/Social_Network_Ads/Social_Network_Ads.csv")

print(df.head())
print(df.info())
print(df.describe())

print(df.isnull().sum())

df.dropna(inplace=True)

plt.figure(figsize=(10, 6))
snb.scatterplot(data=df, x='EstimatedSalary', y='Age', hue='Purchased',
palette='viridis', alpha=0.7)
plt.title('Scatter Plot of Age vs Estimated Salary')
plt.xlabel('Estimated Salary')
plt.ylabel('Age')
plt.legend(title='Purchased', loc='upper left', labels=['No', 'Yes'])
plt.show()

# Encoding categorical data
df['Gender'] = pd.factorize(df['Gender'])[0] # Male=0, Female=1
# Features and target variable
X = df[['Gender', 'Age', 'EstimatedSalary']]
y = df['Purchased']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Train the SVM model
svm_classifier_linear = SVC(kernel='linear')
svm_classifier_linear.fit(X_train, y_train)

y_pred_linear = svm_classifier_linear.predict(X_test)

```

```

# confusion matrix
cm = confusion_matrix(y_test, y_pred_linear)
print("Confusion Matrix for linear kernel:\n", cm)

# metrics
accuracy_linear = accuracy_score(y_test, y_pred_linear)
recall_linear = recall_score(y_test, y_pred_linear)
precision_linear = precision_score(y_test, y_pred_linear)
f1_linear = f1_score(y_test, y_pred_linear)

print("For linear kernel")
print(f'Accuracy: {accuracy_linear:.2f}')
print(f'Recall: {recall_linear:.2f}')
print(f'Precision: {precision_linear:.2f}')
print(f'F1 Score: {f1_linear:.2f}')


```

```

PS E:\College\SEM-VII\ML\PRACS> & "C:/Users/Om Shete/AppData/Local/Programs/Python/Python311/
py
  User ID  Gender  Age  EstimatedSalary  Purchased
0  15624510   Male   19           19000           0
1  15810944   Male   35           20000           0
2  15668575  Female   26           43000           0
3  15603246  Female   27           57000           0
4  15804002   Male   19           76000           0
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   User ID         400 non-null   int64
1   Gender          400 non-null   object
2   Age             400 non-null   int64
3   EstimatedSalary 400 non-null   int64
4   Purchased       400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB

```

```
count      User ID      Age  EstimatedSalary  Purchased
mean      1.569154e+07  37.655000    69742.500000    0.357500
std       7.165832e+04  10.482877    34096.960282    0.479864
min       1.556669e+07  18.000000    15000.000000    0.000000
25%       1.562676e+07  29.750000    43000.000000    0.000000
50%       1.569434e+07  37.000000    70000.000000    0.000000
75%       1.575036e+07  46.000000    88000.000000    1.000000
max       1.581524e+07  60.000000   150000.000000    1.000000
User ID    0
Gender      0
Age         0
EstimatedSalary  0
Purchased   0
dtype: int64
Confusion Matrix for linear kernel:
[[48  4]
 [ 9 19]]
For linear kernel
Accuracy: 0.84
Recall: 0.68
Precision: 0.83
F1 Score: 0.75
```



```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as snb
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import confusion_matrix, accuracy_score,
recall_score, precision_score, f1_score

df = pd.read_csv("dataset/Social_Network_Ads/Social_Network_Ads.csv")

print(df.head())
print(df.info())
print(df.describe())

print(df.isnull().sum())
df.dropna(inplace=True)

plt.figure(figsize=(10, 6))
snb.scatterplot(data=df, x='EstimatedSalary', y='Age',
hue='Purchased', palette='viridis', alpha=0.7)
plt.title('Scatter Plot of Age vs Estimated Salary')
plt.xlabel('Estimated Salary')
plt.ylabel('Age')
plt.legend(title='Purchased', loc='upper left', labels=['No', 'Yes'])
plt.show()

# Encoding categorical data
df['Gender'] = pd.factorize(df['Gender'])[0] # Male=0, Female=1
# Features and target variable
X = df[['Gender', 'Age', 'EstimatedSalary']]
y = df['Purchased']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

svm_model = SVC(kernel='rbf')
svm_model.fit(X_train, y_train)

y_pred = svm_model.predict(X_test)

cm = confusion_matrix(y_test, y_pred)

```

```

print("Confusion matrix: \n", cm)

accuracy= accuracy_score(y_test, y_pred)
recall= recall_score(y_test, y_pred)
precision= precision_score(y_test, y_pred)
f1= f1_score(y_test, y_pred)

print(f'Accuracy: {accuracy:.2f}')
print(f'Recall: {recall:.2f}')
print(f'Precision: {precision:.2f}')
print(f'F1 Score: {f1:.2f}')

```

```


PS E:\College\SEM-VII\ML\PRACS> & "C:/Users/Om Shete/AppData/Local/Programs/Python/Python311/pyth
py
   User ID  Gender  Age  EstimatedSalary  Purchased
0  15624510   Male   19             19000           0
1  15810944   Male   35             20000           0
2  15668575  Female   26             43000           0
3  15603246  Female   27             57000           0
4  15804002   Male   19             76000           0
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   User ID              400 non-null    int64
 1   Gender               400 non-null    object
 2   Age                 400 non-null    int64
 3   EstimatedSalary      400 non-null    int64
 4   Purchased           400 non-null    int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
None

```

```

count    4.000000e+02    400.000000    400.000000    400.000000
mean     1.569154e+07    37.655000    69742.500000    0.357500
std      7.165832e+04    10.482877    34096.960282    0.479864
min      1.556669e+07    18.000000    15000.000000    0.000000
25%      1.562676e+07    29.750000    43000.000000    0.000000
50%      1.569434e+07    37.000000    70000.000000    0.000000
75%      1.575036e+07    46.000000    88000.000000    1.000000
max      1.581524e+07    60.000000    150000.000000    1.000000
User ID      0
Gender       0
Age          0
EstimatedSalary  0
Purchased    0
dtype: int64
Confusion matrix:
[[49  3]
 [18 10]]
Accuracy: 0.74
Recall: 0.36
Precision: 0.77
F1 Score: 0.49

```




```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score,
recall_score, precision_score, f1_score
from sklearn.svm import SVC
import matplotlib.pyplot as plt
import seaborn as snb

df = pd.read_csv("dataset/Iris/Iris.csv")

print(df.head())
print(df.info())
print(df.describe())

df['Species'] = pd.factorize(df['Species'])[0]

print(df.isnull().sum())
df.dropna(inplace=True)

X = df.drop(['Id', 'Species'], axis=1)
y = df['Species']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

svm_classifier_linear = SVC(kernel='linear')
svm_classifier_linear.fit(X_train, y_train)

y_pred_linear = svm_classifier_linear.predict(X_test)

cm = confusion_matrix(y_test, y_pred_linear)
print("Confusion Matrix for linear kernel:\n", cm)

accuracy_linear = accuracy_score(y_test, y_pred_linear)
recall_linear = recall_score(y_test, y_pred_linear, average='weighted')
precision_linear = precision_score(y_test, y_pred_linear,
average='weighted')
f1_linear = f1_score(y_test, y_pred_linear, average='weighted')

print(f'Accuracy: {accuracy_linear:.2f}')
print(f'Recall: {recall_linear:.2f}')

```

```
print(f'Precision: {precision_linear:.2f}')
print(f'F1 Score: {f1_linear:.2f}')
```

```
PS E:\College\SEM-VII\ML\PRACS> & "C:/Users/Om Shete/AppData/Local/Programs/Python/Python311/python.exe"
py
   Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  Species
0    1             5.1             3.5             1.4             0.2  Iris-setosa
1    2             4.9             3.0             1.4             0.2  Iris-setosa
2    3             4.7             3.2             1.3             0.2  Iris-setosa
3    4             4.6             3.1             1.5             0.2  Iris-setosa
4    5             5.0             3.6             1.4             0.2  Iris-setosa
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0    Id              150 non-null    int64
1    SepalLengthCm   150 non-null    float64
2    SepalWidthCm    150 non-null    float64
3    PetalLengthCm   150 non-null    float64
4    PetalWidthCm    150 non-null    float64
5    Species         150 non-null    object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
None
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	75.500000	5.843333	3.054000	3.758667	1.198667
std	43.445368	0.828066	0.433594	1.764420	0.763161
min	1.000000	4.300000	2.000000	1.000000	0.100000
25%	38.250000	5.100000	2.800000	1.600000	0.300000
50%	75.500000	5.800000	3.000000	4.350000	1.300000
75%	112.750000	6.400000	3.300000	5.100000	1.800000
max	150.000000	7.900000	4.400000	6.900000	2.500000
Id	0				
SepalLengthCm	0				
SepalWidthCm	0				
PetalLengthCm	0				
PetalWidthCm	0				
Species	0				
dtype:	int64				

```
Confusion Matrix for linear kernel:
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
Accuracy: 1.00
Recall: 1.00
Precision: 1.00
F1 Score: 1.00
```

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score,
recall_score, precision_score, f1_score
from sklearn.svm import SVC
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv("dataset/Iris/Iris.csv")

print(df.head())
print(df.info())
print(df.describe())

df['Species'] = pd.factorize(df['Species'])[0]

print(df.isnull().sum())
df.dropna(inplace=True)

X = df.drop(['Id', 'Species'], axis=1)
y = df['Species']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

svm_classifier_rbf = SVC(kernel='rbf')
svm_classifier_rbf.fit(X_train, y_train)

y_pred_rbf = svm_classifier_rbf.predict(X_test)

cm = confusion_matrix(y_test, y_pred_rbf)
print("Confusion Matrix for rbf kernel:\n", cm)

accuracy_rbf = accuracy_score(y_test, y_pred_rbf)
recall_rbf = recall_score(y_test, y_pred_rbf, average='weighted')
precision_rbf = precision_score(y_test, y_pred_rbf, average='weighted')
f1_rbf = f1_score(y_test, y_pred_rbf, average='weighted')

print(f'Accuracy: {accuracy_rbf:.2f}')
print(f'Recall: {recall_rbf:.2f}')

```

```
print(f'Precision: {precision_rbf:.2f}')
print(f'F1 Score: {f1_rbf:.2f}')
```

```
PS E:\College\SEM-VII\ML\PRACS> & "C:/Users/Om Shete/AppData/Local/Programs/Python/Python311/python.exe" e:/College/SEM-
py
   Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  Species
0    1             5.1           3.5           1.4           0.2  Iris-setosa
1    2             4.9           3.0           1.4           0.2  Iris-setosa
2    3             4.7           3.2           1.3           0.2  Iris-setosa
3    4             4.6           3.1           1.5           0.2  Iris-setosa
4    5             5.0           3.6           1.4           0.2  Iris-setosa

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  ---
0    Id              150 non-null    int64
1    SepalLengthCm    150 non-null    float64
2    SepalWidthCm     150 non-null    float64
3    PetalLengthCm    150 non-null    float64
4    PetalWidthCm     150 non-null    float64
5    Species          150 non-null    object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
None
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	75.500000	5.843333	3.054000	3.758667	1.198667
std	43.445368	0.828066	0.433594	1.764420	0.763161
min	1.000000	4.300000	2.000000	1.000000	0.100000
25%	38.250000	5.100000	2.800000	1.600000	0.300000
50%	75.500000	5.800000	3.000000	4.350000	1.300000
75%	112.750000	6.400000	3.300000	5.100000	1.800000
max	150.000000	7.900000	4.400000	6.900000	2.500000
Id	0				
SepalLengthCm	0				
SepalWidthCm	0				
PetalLengthCm	0				
PetalWidthCm	0				
Species	0				

```
dtype: int64
Confusion Matrix for rbf kernel:
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
Accuracy: 1.00
Recall: 1.00
Precision: 1.00
F1 Score: 1.00
```

```

#35. The Iris data set contains 3 classes of 50 instances each, where
each class refers to a type of iris plant.
# i. Understand the Dataset & cleanup (if required).
# ii. Build a decision tree classifier to predict the iris plant
category? (Use GINI INDEX criteria, use
# max_depth=4, min_samples_split=2)
# iii. Evaluate the model using Accuracy.

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import plot_tree
from sklearn.metrics import accuracy_score, confusion_matrix, log_loss

df = pd.read_csv("dataset/Iris/Iris.csv")

print(df.sample(5))
print(df.info())

df['Species'] = pd.factorize(df['Species'])[0]

print(df.isnull().sum())

df.dropna(inplace=True)

X = df.drop(['Id', 'Species'], axis=1)
y = df['Species']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

dt_classifier = DecisionTreeClassifier(criterion="gini", max_depth=4,
min_samples_split=2, random_state=42)
dt_classifier.fit(X_train, y_train)

y_pred = dt_classifier.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy Score: {accuracy}")

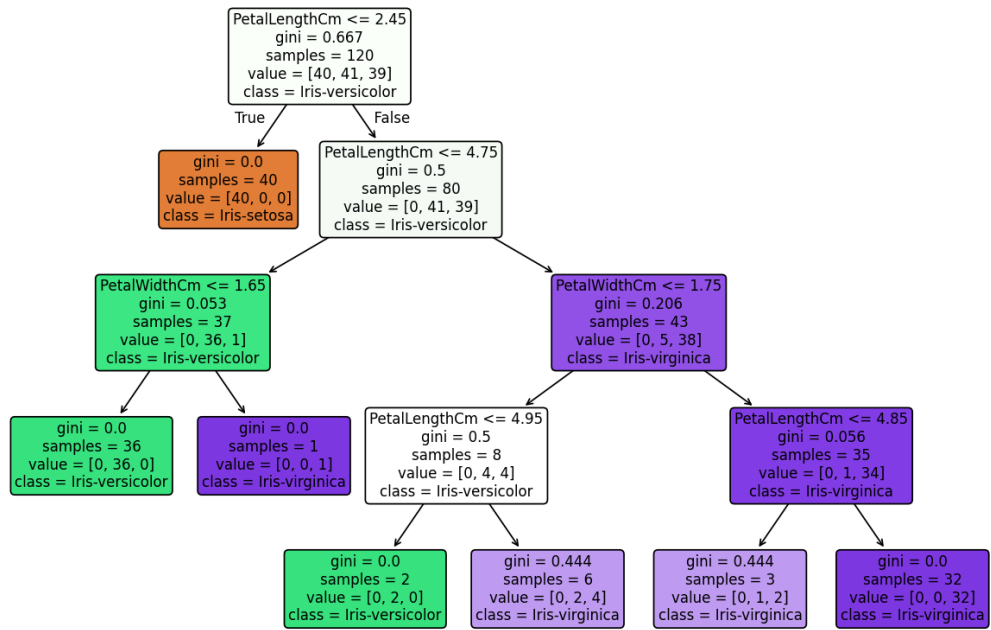
```

```
plt.figure(figsize=(12,8))
plot_tree(dt_classifier, filled=True, feature_names=X.columns,
class_names=["Iris-setosa","Iris-versicolor","Iris-virginica"],
rounded=True)
plt.title("Decision Tree Visualization")
plt.show()
```

```
PS E:\College\SEM-VII\ML\PRACS> & "C:/Users/Om Shete/AppData/Local/Programs/Python/Python311/pyth
    Id SepalLengthCm SepalWidthCm PetalLengthCm PetalWidthCm Species
116 117          6.5          3.0          5.5          1.8 Iris-virginica
66  67          5.6          3.0          4.5          1.5 Iris-versicolor
15  16          5.7          4.4          1.5          0.4 Iris-setosa
130 131          7.4          2.8          6.1          1.9 Iris-virginica
61  62          5.9          3.0          4.2          1.5 Iris-versicolor
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Id              150 non-null   int64
1   SepalLengthCm   150 non-null   float64
2   SepalWidthCm    150 non-null   float64
3   PetalLengthCm   150 non-null   float64
4   PetalWidthCm    150 non-null   float64
5   Species         150 non-null   object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
None
Id              0
SepalLengthCm   0
```

```
Id              0
SepalLengthCm   0
SepalWidthCm    0
PetalLengthCm   0
PetalWidthCm    0
Species         0
dtype: int64
Accuracy Score: 1.0
```

Decision Tree Visualization



```

# 36. The Iris data set contains 3 classes of 50 instances each, where
each class refers to a type of iris plant.
# i. Understand the Dataset & cleanup (if required).
# ii. Build a decision tree classifier to predict the iris plant
category. (Use Entropy criteria, use
# max_depth=4, min_samples_split=2)
# iii. Evaluate the model using Accuracy.

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import plot_tree
from sklearn.metrics import accuracy_score, confusion_matrix, log_loss

df = pd.read_csv("dataset/Iris/Iris.csv")

print(df.sample(5))
print(df.info())

df['Species'] = pd.factorize(df['Species'])[0]

print(df.isnull().sum())

df.dropna(inplace=True)

X = df.drop(['Id', 'Species'], axis=1)
y = df['Species']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

dt_classifier = DecisionTreeClassifier(criterion="entropy",
max_depth=4, min_samples_split=2, random_state=42)
dt_classifier.fit(X_train, y_train)

y_pred = dt_classifier.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy Score: {accuracy}")

```

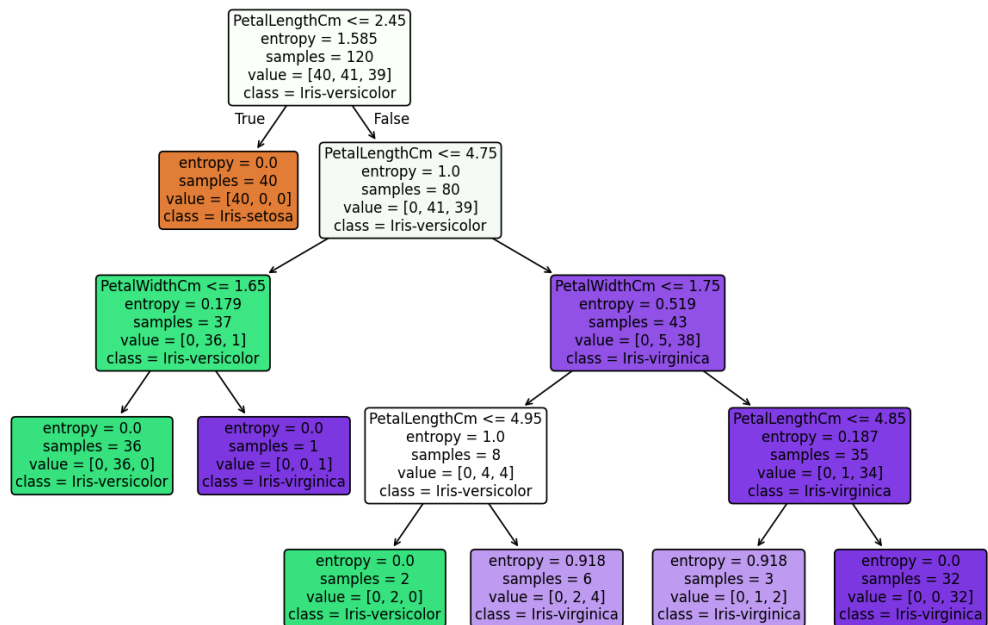


```
plt.figure(figsize=(12,8))
plot_tree(dt_classifier, filled=True, feature_names=X.columns,
class_names=["Iris-setosa","Iris-versicolor","Iris-virginica"],
rounded=True)
plt.title("Decision Tree Visualization")
plt.show()
```

```
PS E:\College\SEM-VII\ML\PRACS> & "C:/Users/Om Shete/AppData/Local/Programs/Python/Python311/python.exe" e:
      Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  Species
115  116           6.4           3.2           5.3           2.3  Iris-virginica
132  133           6.4           2.8           5.6           2.2  Iris-virginica
 94   95           5.6           2.7           4.2           1.3  Iris-versicolor
 28   29           5.2           3.4           1.4           0.2   Iris-setosa
 57   58           4.9           2.4           3.3           1.0  Iris-versicolor
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Id              150 non-null   int64
 1   SepalLengthCm   150 non-null   float64
 2   SepalWidthCm    150 non-null   float64
 3   PetalLengthCm   150 non-null   float64
 4   PetalWidthCm    150 non-null   float64
 5   Species         150 non-null   object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
None
Id              0
SepalLengthCm   0
SepalWidthCm    0
PetalLengthCm   0
```

```
Id              0
SepalLengthCm   0
SepalWidthCm    0
PetalLengthCm   0
PetalWidthCm    0
Species         0
dtype: int64
Accuracy Score: 1.0
```

Decision Tree Visualization



```

#37. The Iris data set contains 3 classes of 50 instances each, where
each class refers to a type of iris plant.
# i. Understand the Dataset & cleanup (if required).
# ii. Build a decision tree classifier to predict the iris plant
category? (Use log loss criteria, use
# max_depth=4, min_samples_split=2)
# iii. Evaluate the model using Accuracy.

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import plot_tree
from sklearn.metrics import accuracy_score, confusion_matrix, log_loss

df = pd.read_csv("dataset/Iris/Iris.csv")

print(df.sample(5))
print(df.info())

df['Species'] = pd.factorize(df['Species'])[0]

print(df.isnull().sum())

df.dropna(inplace=True)

X = df.drop(['Id', 'Species'], axis=1)
y = df['Species']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

dt_classifier = DecisionTreeClassifier(criterion="log_loss",
max_depth=4, min_samples_split=2, random_state=42)
dt_classifier.fit(X_train, y_train)

y_pred = dt_classifier.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy Score: {accuracy}")

```

```
plt.figure(figsize=(12,8))
plot_tree(dt_classifier, filled=True, feature_names=X.columns,
class_names=["Iris-setosa","Iris-versicolor","Iris-virginica"],
rounded=True)
plt.title("Decision Tree Visualization")
plt.show()
```

```
PS E:\College\SEM-VII\ML\PRACS> & "C:/Users/Om Shete/AppData/Local/Programs/Python/Python311/python.exe" e:/C

```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
64	65	5.6	2.9	3.6	1.3	Iris-versicolor
52	53	6.9	3.1	4.9	1.5	Iris-versicolor
125	126	7.2	3.2	6.0	1.8	Iris-virginica
98	99	5.1	2.5	3.0	1.1	Iris-versicolor
127	128	6.1	3.0	4.9	1.8	Iris-virginica

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  ---
0    Id              150 non-null    int64
1    SepalLengthCm   150 non-null    float64
2    SepalWidthCm    150 non-null    float64
3    PetalLengthCm   150 non-null    float64
4    PetalWidthCm    150 non-null    float64
5    Species         150 non-null    object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB

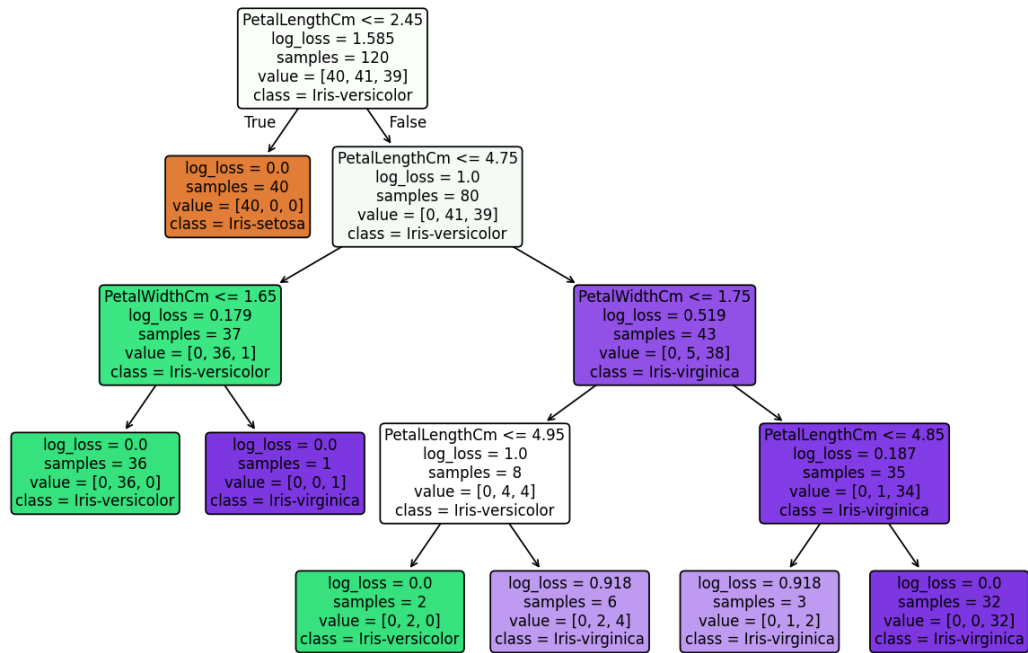
```

```

None
Id              0
SepalLengthCm  0
SepalWidthCm   0
PetalLengthCm  0
PetalWidthCm   0
Species        0
dtype: int64
Accuracy Score: 1.0

```

Decision Tree Visualization



```

# The Iris data set contains 3 classes of 50 instances each, where each
class refers to a type of iris plant.
# i. Understand the Dataset & cleanup (if required).
# ii. Build a logistic Regression classifier to predict the iris plant
category?
# iii. Evaluate the model using Accuracy

import pandas as pd
from sklearn.model_selection import train_test_split
from LogisticRegression import LogisticRegression
from sklearn.metrics import accuracy_score

df = pd.read_csv("dataset/Iris/Iris.csv")

print(df.sample(5))
print(df.info())

df['Species'] = pd.factorize(df['Species'])[0]
print(df.head())

print(df.isnull().sum())
df.dropna(inplace=True)

print(df['Species'].value_counts())

X = df.drop(columns=['Id', 'Species'])
y = df['Species']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

logistic_model = LogisticRegression()
logistic_model.fit(X_train, y_train)

y_pred = logistic_model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Logistic Regression Accuracy: {accuracy:.2f}")

```

```

PS E:\College\SEM-VII\ML\PRACS> & "C:/Users/Om Shete/AppData/Local/Programs/Python/Python311/python.exe" e:/College/3
    Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  Species
143  144            6.8            3.2            5.9            2.3  Iris-virginica
65   66            6.7            3.1            4.4            1.4  Iris-versicolor
11   12            4.8            3.4            1.6            0.2  Iris-setosa
120  121            6.9            3.2            5.7            2.3  Iris-virginica
74   75            6.4            2.9            4.3            1.3  Iris-versicolor
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Id              150 non-null    int64
1   SepalLengthCm   150 non-null    float64
2   SepalWidthCm    150 non-null    float64
3   PetalLengthCm   150 non-null    float64
4   PetalWidthCm    150 non-null    float64
5   Species         150 non-null    object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
None

```

```

    Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  Species
0   1            5.1            3.5            1.4            0.2      0
1   2            4.9            3.0            1.4            0.2      0
2   3            4.7            3.2            1.3            0.2      0
3   4            4.6            3.1            1.5            0.2      0
4   5            5.0            3.6            1.4            0.2      0
Id              0
SepalLengthCm   0
SepalWidthCm    0
PetalLengthCm   0
PetalWidthCm    0
Species         0
dtype: int64
Species
0    50
1    50
2    50
Name: count, dtype: int64
Logistic Regression Accuracy: 0.30

```

```

# 39. The Iris data set contains 3 classes of 50 instances each, where
each class refers to a type of iris plant.
# i. Understand the Dataset & cleanup (if required).
# ii. Build a Bagging Classifier model to predict the iris plant
category?
# iii. Evaluate the model using Accuracy.

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import BaggingClassifier
from sklearn.metrics import accuracy_score

df = pd.read_csv("dataset/Iris/Iris.csv")

print(df.sample(5))
print(df.info())

df['Species'] = pd.factorize(df['Species'])[0]
print(df.head())

print(df.isnull().sum())
df.dropna(inplace=True)

print(df['Species'].value_counts())

X = df.drop(columns=['Id', 'Species'])
y = df['Species']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

bagging_model = BaggingClassifier()
bagging_model.fit(X_train, y_train)

y_pred = bagging_model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Bagging Classifier Accuracy: {accuracy:.2f}")

```



```

PS E:\College\SEM-VII\ML\PRACS> & "C:/Users/Om Shete/AppData/Local/Programs/Python/Python311/python.exe" e:/Col
    Id SepalLengthCm SepalWidthCm PetalLengthCm PetalWidthCm Species
75    76           6.6           3.0           4.4           1.4 Iris-versicolor
8      9           4.4           2.9           1.4           0.2 Iris-setosa
135   136           7.7           3.0           6.1           2.3 Iris-virginica
43    44           5.0           3.5           1.6           0.6 Iris-setosa
66    67           5.6           3.0           4.5           1.5 Iris-versicolor

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
 #   Column          Non-Null Count  Dtype  
---  --
 0   Id              150 non-null   int64  
 1   SepalLengthCm   150 non-null   float64
 2   SepalWidthCm    150 non-null   float64
 3   PetalLengthCm   150 non-null   float64
 4   PetalWidthCm    150 non-null   float64
 5   Species         150 non-null   object  
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
None

```

```

    Id SepalLengthCm SepalWidthCm PetalLengthCm PetalWidthCm Species
0     1           5.1           3.5           1.4           0.2      0
1     2           4.9           3.0           1.4           0.2      0
2     3           4.7           3.2           1.3           0.2      0
3     4           4.6           3.1           1.5           0.2      0
4     5           5.0           3.6           1.4           0.2      0
Id              0
SepalLengthCm   0
SepalWidthCm    0
PetalLengthCm   0
PetalWidthCm    0
Species         0
dtype: int64
Species
0     50
1     50
2     50
Name: count, dtype: int64
Bagging Classifier Accuracy: 1.00

```

```

# 40. The Iris data set contains 3 classes of 50 instances each, where
each class refers to a type of iris plant.
# i. Understand the Dataset & cleanup (if required).
# ii. Build a Random Forest Classifier model to predict the iris plant
category?
# iii. Evaluate the model using Accuracy

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

df = pd.read_csv("dataset/Iris/Iris.csv")

print(df.sample(5))
print(df.info())

df['Species'] = pd.factorize(df['Species'])[0]
print(df.head())

print(df.isnull().sum())
df.dropna(inplace=True)

print(df['Species'].value_counts())

X = df.drop(columns=['Id', 'Species'])
y = df['Species']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

random_forest_model = RandomForestClassifier()
random_forest_model.fit(X_train, y_train)

y_pred = random_forest_model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Random Forest Classifier Accuracy: {accuracy:.2f}")

```

```

PS E:\College\SEM-VII\ML\PRACS> & "C:/Users/Om Shete/AppData/Local/Programs/Python/Python311/python.exe" e:/College/SEM-VII/ML/PRACS/q40.py
   Id  SepallengthCm  SepalWidthCm  PetallengthCm  PetalWidthCm  Species
10  11             5.4           3.7           1.5           0.2      Iris-setosa
52  53             6.9           3.1           4.9           1.5  Iris-versicolor
108 109             6.7           2.5           5.8           1.8  Iris-virginica
107 108             7.3           2.9           6.3           1.8  Iris-virginica
116 117             6.5           3.0           5.5           1.8  Iris-virginica

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Id              150 non-null    int64
 1   SepallengthCm   150 non-null    float64
 2   SepalWidthCm    150 non-null    float64
 3   PetallengthCm   150 non-null    float64
 4   PetalWidthCm    150 non-null    float64
 5   Species         150 non-null    object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
None

```

```

   Id  SepallengthCm  SepalWidthCm  PetallengthCm  PetalWidthCm  Species
0    1             5.1           3.5           1.4           0.2      0
1    2             4.9           3.0           1.4           0.2      0
2    3             4.7           3.2           1.3           0.2      0
3    4             4.6           3.1           1.5           0.2      0
4    5             5.0           3.6           1.4           0.2      0
Id
SepallengthCm
SepalWidthCm
PetallengthCm
PetalWidthCm
Species
dtype: int64
Species
0    50
1    50
2    50
Name: count, dtype: int64
Random Forest Classifier Accuracy: 1.00

```

```

# 41. The Iris data set contains 3 classes of 50 instances each, where
each class refers to a type of iris plant.
# i. Understand the Dataset & cleanup (if required).
# ii. Build a Gradient Boost Classifier model to predict the iris plant
category?
# iii. Evaluate the model using Accuracy

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import accuracy_score

df = pd.read_csv("dataset/Iris/Iris.csv")

print(df.sample(5))
print(df.info())

df['Species'] = pd.factorize(df['Species'])[0]
print(df.head())

print(df.isnull().sum())
df.dropna(inplace=True)

print(df['Species'].value_counts())

X = df.drop(columns=['Id', 'Species'])
y = df['Species']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

gradient_boosting_model = GradientBoostingClassifier()
gradient_boosting_model.fit(X_train, y_train)

y_pred = gradient_boosting_model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Gradient Boosting Classifier Accuracy: {accuracy:.2f}")

```

```

PS E:\College\SEM-VII\ML\PRACS> & "C:/Users/Om Shete/AppData/Local/Programs/Python/Python311/python.exe" e:/College/SEM-V
    Id SepalLengthCm SepalWidthCm PetalLengthCm PetalWidthCm Species
130 131           7.4           2.8           6.1           1.9 Iris-virginica
106 107           4.9           2.5           4.5           1.7 Iris-virginica
 71  72           6.1           2.8           4.0           1.3 Iris-versicolor
 63  64           6.1           2.9           4.7           1.4 Iris-versicolor
 90  91           5.5           2.6           4.4           1.2 Iris-versicolor

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
#   Column             Non-Null Count  Dtype
---  -
0   Id                  150 non-null   int64
1   SepalLengthCm       150 non-null   float64
2   SepalWidthCm        150 non-null   float64
3   PetalLengthCm       150 non-null   float64
4   PetalWidthCm        150 non-null   float64
5   Species             150 non-null   object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
None

```

```

    Id SepalLengthCm SepalWidthCm PetalLengthCm PetalWidthCm Species
0    1           5.1           3.5           1.4           0.2      0
1    2           4.9           3.0           1.4           0.2      0
2    3           4.7           3.2           1.3           0.2      0
3    4           4.6           3.1           1.5           0.2      0
4    5           5.0           3.6           1.4           0.2      0
Id
SepalLengthCm
SepalWidthCm
PetalLengthCm
PetalWidthCm
Species
dtype: int64
Species
0    50
1    50
2    50
Name: count, dtype: int64
Gradient Boosting Classifier Accuracy: 1.00

```

```

# 42. The Iris data set contains 3 classes of 50 instances each, where
each class refers to a type of iris plant.
# i. Understand the Dataset & cleanup (if required).
# ii. Build a AdaBoost Classifier model to predict the iris plant
category?
# iii. Evaluate the model using Accuracy

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import AdaBoostClassifier
from sklearn.metrics import accuracy_score

df = pd.read_csv("dataset/Iris/Iris.csv")

print(df.sample(5))
print(df.info())

df['Species'] = pd.factorize(df['Species'])[0]
print(df.head())

print(df.isnull().sum())
df.dropna(inplace=True)

print(df['Species'].value_counts())

X = df.drop(columns=['Id', 'Species'])
y = df['Species']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

ada_boost_model = AdaBoostClassifier()
ada_boost_model.fit(X_train, y_train)

y_pred = ada_boost_model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Ada Boost Classifier Accuracy: {accuracy:.2f}")

```

```

      Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  Species
108  109             6.7             2.5             5.8             1.8  Iris-virginica
137  138             6.4             3.1             5.5             1.8  Iris-virginica
145  146             6.7             3.0             5.2             2.3  Iris-virginica
115  116             6.4             3.2             5.3             2.3  Iris-virginica
   4     5             5.0             3.6             1.4             0.2  Iris-setosa

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   Id                    150 non-null    int64  
1   SepalLengthCm         150 non-null    float64
2   SepalWidthCm          150 non-null    float64
3   PetalLengthCm         150 non-null    float64
4   PetalWidthCm          150 non-null    float64
5   Species               150 non-null    object  
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
None

```

```

      Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  Species
0     1             5.1             3.5             1.4             0.2         0
1     2             4.9             3.0             1.4             0.2         0
2     3             4.7             3.2             1.3             0.2         0
3     4             4.6             3.1             1.5             0.2         0
4     5             5.0             3.6             1.4             0.2         0
Id
SepalLengthCm
SepalWidthCm
PetalLengthCm
PetalWidthCm
Species
dtype: int64
Species
0     50
1     50
2     50
Name: count, dtype: int64
C:\Users\Om Shete\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:17: DeprecationWarning: The
ME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.
warnings.warn(
Ada Boost Classifier Accuracy: 1.00

```

```

import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.ensemble import RandomForestClassifier
from sklearn.decomposition import PCA

df = pd.read_csv("dataset/Iris/Iris.csv")

print(df.sample(5))
print(df.info())

df['Species'] = pd.factorize(df['Species'])[0]

print(df.sample(5))

print(df.isnull().sum())
df.dropna(inplace=True)

X = df.drop(columns=['Id', 'Species'])
y = df['Species']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Without PCA

random_forest_model = RandomForestClassifier(random_state=42)
random_forest_model.fit(X_train, y_train)

y_pred = random_forest_model.predict(X_test)

accuracy_without_pca = accuracy_score(y_test, y_pred)

print("Confusion Matrix without PCA: \n", confusion_matrix(y_test,
y_pred))
print("Accuracy without PCA: \n", accuracy_without_pca)

# With PCA

pca = PCA(n_components=2)

```



```
X_pca = pca.fit_transform(X)

X_train_pca, X_test_pca, y_train_pca, y_test_pca =
train_test_split(X_pca, y, test_size=0.2, random_state=42)

rf_classifier_pca = RandomForestClassifier(random_state=42)
rf_classifier_pca.fit(X_train_pca, y_train_pca)

y_pred_pca = rf_classifier_pca.predict(X_test_pca)

accuracy_with_pca = accuracy_score(y_test_pca, y_pred_pca)
print("Confusion Matrix:\n", confusion_matrix(y_test_pca, y_pred_pca))
print(f"Accuracy with PCA: {accuracy_with_pca:.2f}")

#

print(f"Accuracy before PCA: {accuracy_without_pca:.2f}")
print(f"Accuracy after PCA: {accuracy_with_pca:.2f}")

#

plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y, cmap='viridis')
plt.xlabel('First Principal Component')
plt.ylabel('Second Principal Component')
plt.title('Iris Dataset after PCA')
plt.colorbar()
plt.show()
```

```
PS E:\College\SEM-VII\ML\PRACS> & "C:/Users/Om Shete/AppData/Local/Programs/Python/Python3
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
77	78	6.7	3.0	5.0	1.7	Iris-versicolor
121	122	5.6	2.8	4.9	2.0	Iris-virginica
74	75	6.4	2.9	4.3	1.3	Iris-versicolor
134	135	6.1	2.6	5.6	1.4	Iris-virginica
8	9	4.4	2.9	1.4	0.2	Iris-setosa

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 150 entries, 0 to 149
```

```
Data columns (total 6 columns):
```

#	Column	Non-Null Count	Dtype
0	Id	150 non-null	int64
1	SepalLengthCm	150 non-null	float64
2	SepalWidthCm	150 non-null	float64
3	PetalLengthCm	150 non-null	float64
4	PetalWidthCm	150 non-null	float64
5	Species	150 non-null	object

123	124	6.3	2.7	4.9	1.8	2
126	127	6.2	2.8	4.8	1.8	2
4	5	5.0	3.6	1.4	0.2	0
110	111	6.5	3.2	5.1	2.0	2
142	143	5.8	2.7	5.1	1.9	2

```
Id          0
SepalLengthCm  0
SepalWidthCm  0
PetalLengthCm  0
PetalWidthCm  0
Species      0
```

```
dtype: int64
```

```
Confusion Matrix without PCA:
```

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

```
Accuracy without PCA:
```

```
1.0
```

```
Confusion Matrix:
```

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

```
Accuracy with PCA: 1.00
```

```
Accuracy before PCA: 1.00
```

```
Accuracy after PCA: 1.00
```

